

A Novel Bidding Strategy for PDAs using MCTS in Continuous Action Spaces

Sanjay Chandekar^{1,2}, Easwar Subramanian¹

1



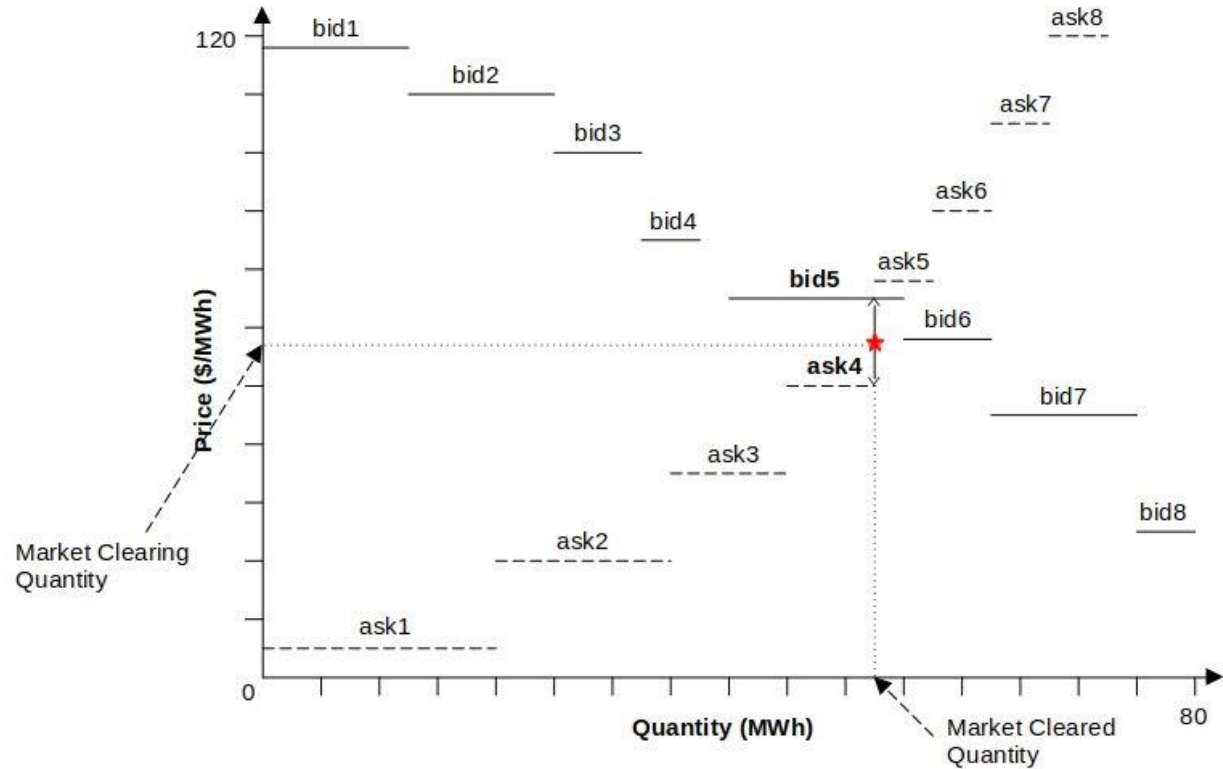
2



Background - PDA


What is Periodic Double Auctions (PDAs)?

Single-shot auction with
k-Double Auction



What is Periodic Double Auctions (PDAs)?

**Periodic Double Auction
(Sequence of auctions)**



Applications: Stock market auction and Energy auction

Challenges in Periodic Double Auctions (PDAs)

- Strategic planning across current and future auctions
- Each decision influences subsequent steps
- Need to be real-time
- Need to account for the updated requirements based on factors like weather

Consequently, the decision-making process for bidding strategies becomes intricate, demanding innovative approaches to navigate the complexities of real-time bidding in PDAs.

**For such sequential
decision-making
problems**

Perfect Fit




**Monte Carlo Tree Search
(MCTS)**


Background - MCTS

Introduction of MCTS


- Monte Carlo Tree Search (MCTS) framework is capable of constructing trees
 - That encompass entire decision-making trajectories
 - Comprehensively capturing process dynamics
- Ability to blend the precision of tree search with the generality of random sampling




Selection
Tree traversed using
tree policy



Expansion
New node added to the
tree (selected using the
tree policy)



Simulation
Rollouts are played
from new node using
default policy



Back-propagation
Final state value is
backpropagated to
parent nodes

Types of MCTS

- Discrete MCTS
- Continuous MCTS
 - Discretization
 - Progressive Widening

Limitation: They do not leverage insights gained from explored actions to perform the exploration of unexplored actions or enhance knowledge about previously explored ones

- Some previous work (mostly for game playing, not directly applicable to other domains)
 - Policy gradient with MCTS

Our Contributions

- We investigate the efficacy of MCTS in continuous action spaces, specifically for the context of placing bids in a periodic double auction.
- Introducing Regression-MCTS (R-MCTS), a novel method designed to navigate the continuous action space of bid prices, harnessing insights gleaned from explored actions, to generalize the action space and accelerating MCTS learning.
- We demonstrate the effectiveness of R-MCTS by evaluating against several state-of-the-art MCTS methods, as well as various PDA bidding strategies, leveraging a simulated PDA environment as our experimental test bed.

Elements of Regression-MCTS (R-MCTS)

- **UCT Selection:** $\operatorname{argmax}_a \left[\bar{v}_a + C \sqrt{\frac{\log \sum_b n_b}{n_a}} \right]$
- **SPW:** $n(v_c)^\alpha \geq |\operatorname{children}(v_c)|$
- **Clearing Probability:**
$$p_{\text{cleared}}(s, \text{action}) = \frac{\sum_{ac \in \text{auction}[s], ac.CP < \text{action}} ac.\text{cleared_amount}}{\sum_{ac \in \text{auction}[s]} ac.\text{cleared_amount}}$$

Regression-MCTS (R-MCTS)

Algorithm 1 R-MCTS(*rem_quant*, *cur_ts*, *delv_ts*)

```
1: bids  $\leftarrow$  [], root  $\leftarrow$  Node() # initialise bids list and root node
2: rem_auctions  $\leftarrow$  delv_ts - cur_ts # number of auctions remaining in a PDA
3: if rem_quant > 0 then
4:   while i in NUMBER_OF_ROLLOUTS do
5:     visited  $\leftarrow$  [], rewards  $\leftarrow$  [], cur  $\leftarrow$  root # initialise lists of visited, rewards
6:     visited  $\leftarrow$  [visited; root]
7:     list_of_sellers  $\leftarrow$  generate_sellers()
8:     list_of_buyers  $\leftarrow$  clone_buyers()
9:     while not cur.is_leaf(rem_quant) do # see algorithm 2
10:      cur  $\leftarrow$  select(rem_quant)
11:      cur.p_cleared  $\leftarrow$  get_pcleared(rem_auctions, cur.action)
12:      cp, cq, rem_quant  $\leftarrow$  perform_auction(cur.action, rem_quant,
        list_of_sellers, list_of_buyers) # clearing the current auction round
13:      rewards  $\leftarrow$  [rewards; cp]
14:      visited  $\leftarrow$  [visited; cur]
15:    end while
16:    cur_cost, cur_quant  $\leftarrow$  cur.simulation(rem_quant, rem_auctions)
17:    root  $\leftarrow$  backpropogation(rewards, visited, cur_cost, cur_quant)
18:  end while
19:  lp  $\leftarrow$  root.best_action()
20:  bids  $\leftarrow$  [bids; Bid(buyer_ID, lp, rem_quant)] # limit order
21: else
22:  bids  $\leftarrow$  [bids; Bid(buyer_ID, NULL, rem_quant)] # market order
23: end if
24: return bids
```

Selection and
Expansion

Simulation
Backpropogation

Best Move

Algorithm 2 $\text{select}(node, rem_quant)$

1: **if** $number_of_visits(node)^\alpha \leq number_of_children(node)$ **then**

2: $A \leftarrow$ actions considered in $node$

3: $action \leftarrow \operatorname{argmax}_{a \in A} \mathbb{E}(v|a) + C \sqrt{\frac{\log \sum_{b \in A} number_of_visits(b)}{number_of_visits(a)}}$ # UCB-select

4: **else**

5: new $action \leftarrow$ child of $node$ by taking an action using $p_cleared$ data

6: $node.children \leftarrow [node.children; action]$ # SPW-select: expanding action space

7: **end if**


8: $new_state = node.action$

9: **return** new_state

SPW
Condition

Variable Risk
Tactic

Variable Risk Tactic



At the start

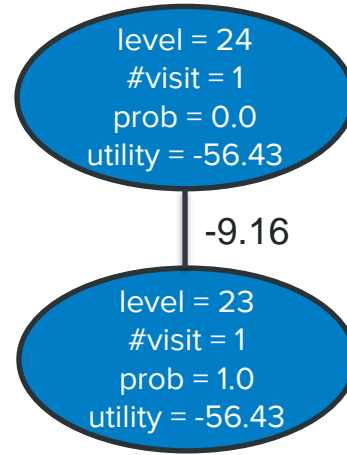
level = 24

#visit = 0


prob = 0.0

utility = 0.0

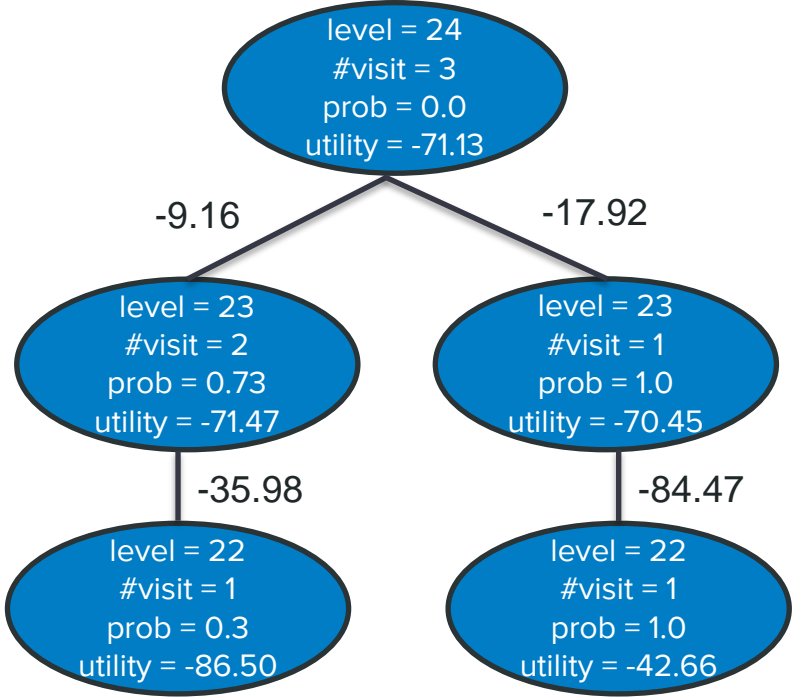
1st Iteration



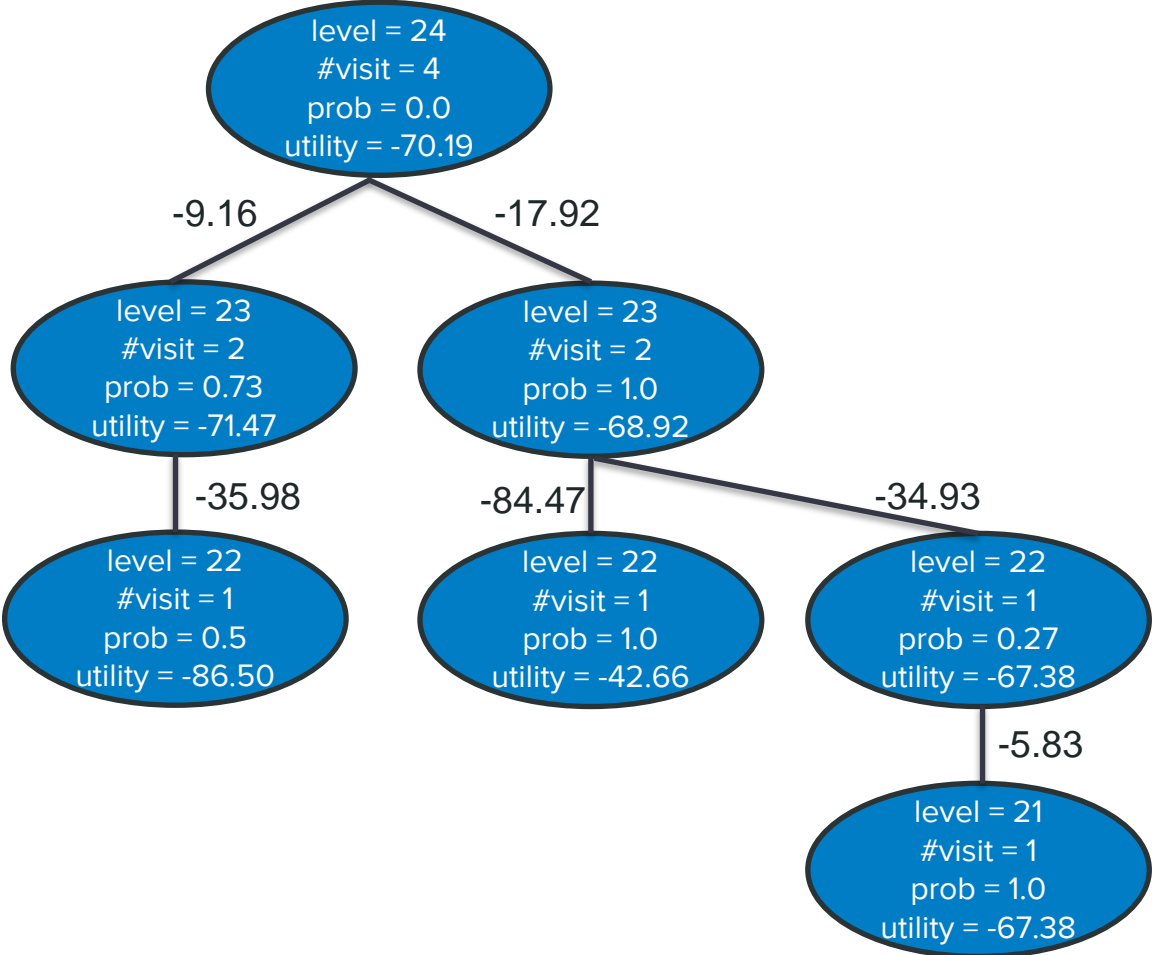
2nd Iteration



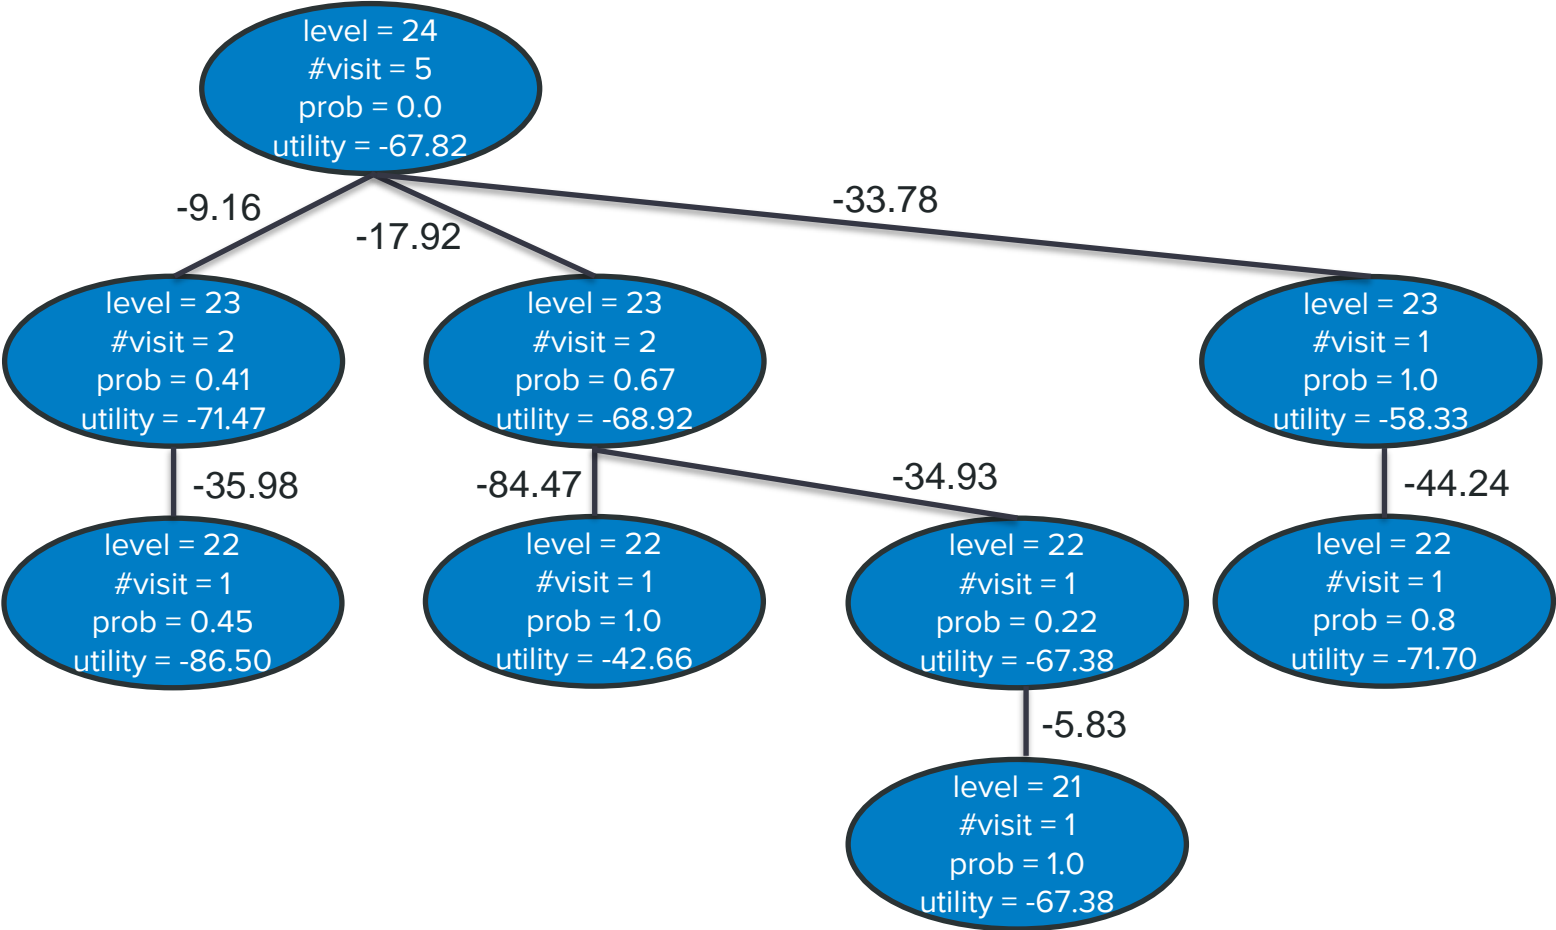
3rd Iteration



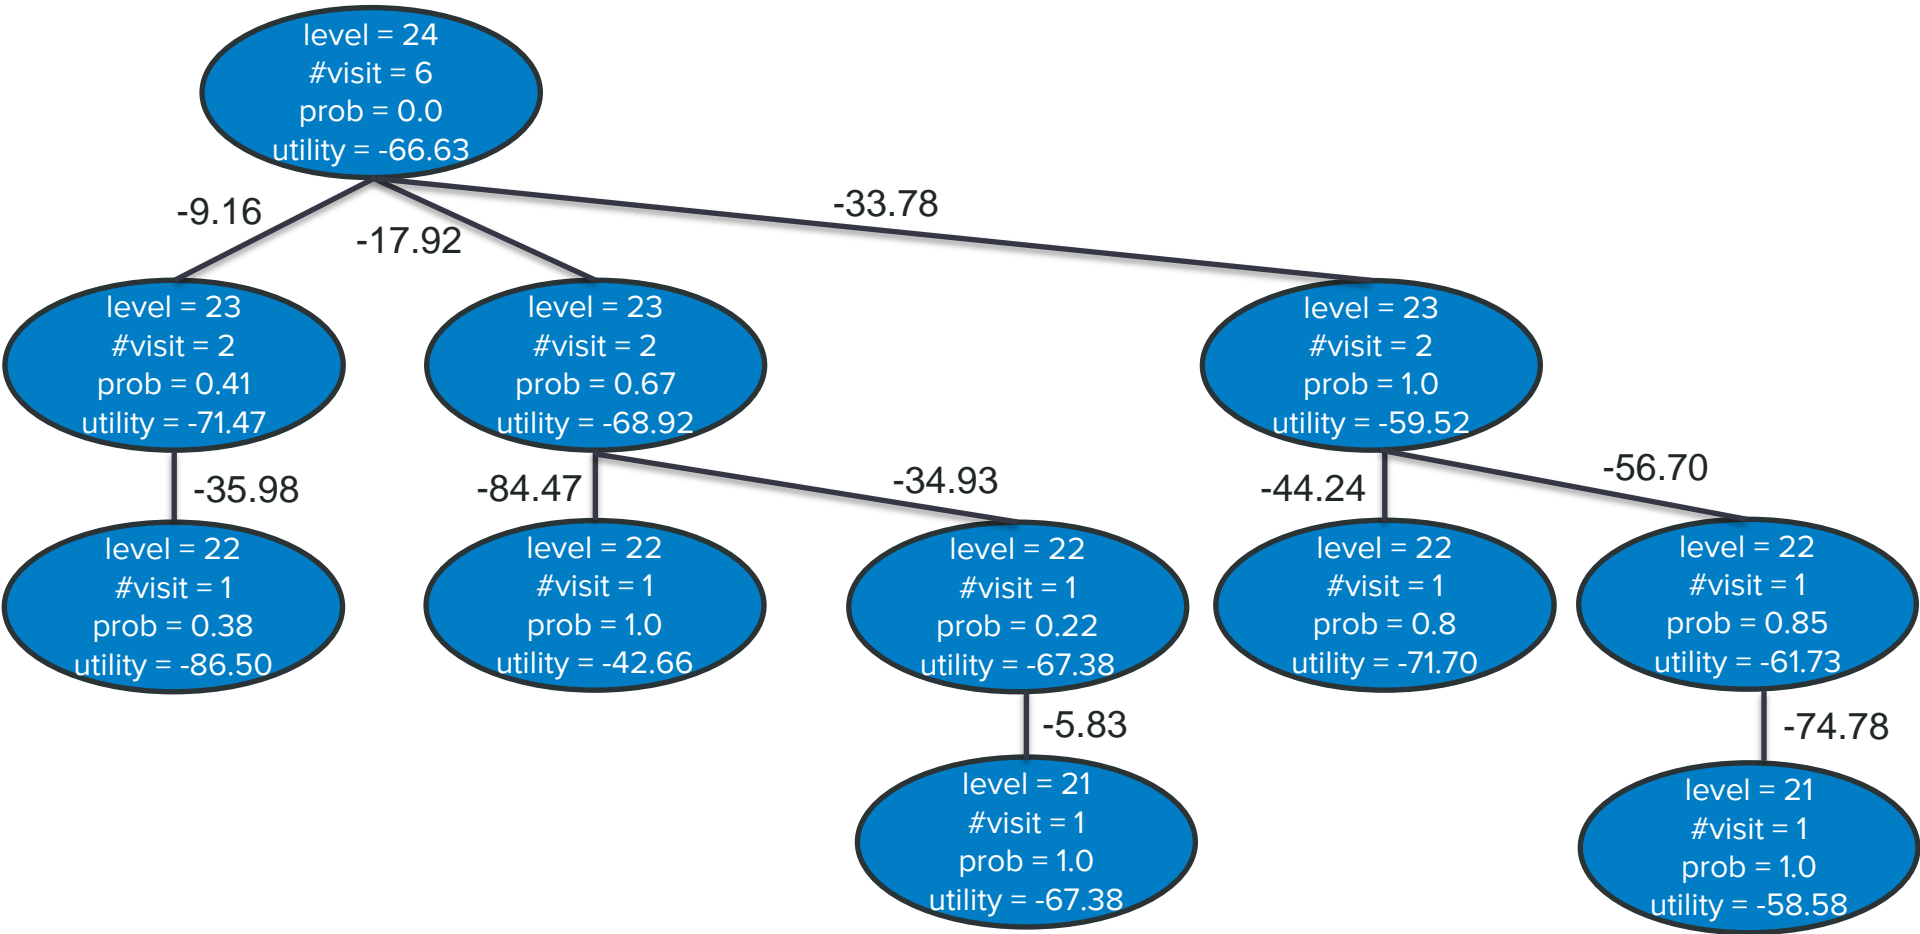
4th Iteration



5th Iteration



5th Iteration



Experiments and Results

Baseline Bidding Strategies for Comparison

MCTS-based Strategies:

MCTS-Vanilla: Discretizes the action space of bid prices. Each action within this space is defined by two multipliers, α_1 and α_2 , applied to the lower and upper bounds of feasible bid prices, respectively.

MCTS-SPW: Initially starts with a discrete action space and dynamically grows its size with the number of rollouts as more bid prices are randomly sampled by balancing between exploration and exploitation.

SPOT: Integrates several heuristic techniques to optimize bid prices and strategically place multiple bids in auctions. Specifically, it calculates the standard deviation of clearing prices σ offline and incorporates an external limit price predictor that provides limit-price μ .

Baseline Bidding Strategies for Comparison

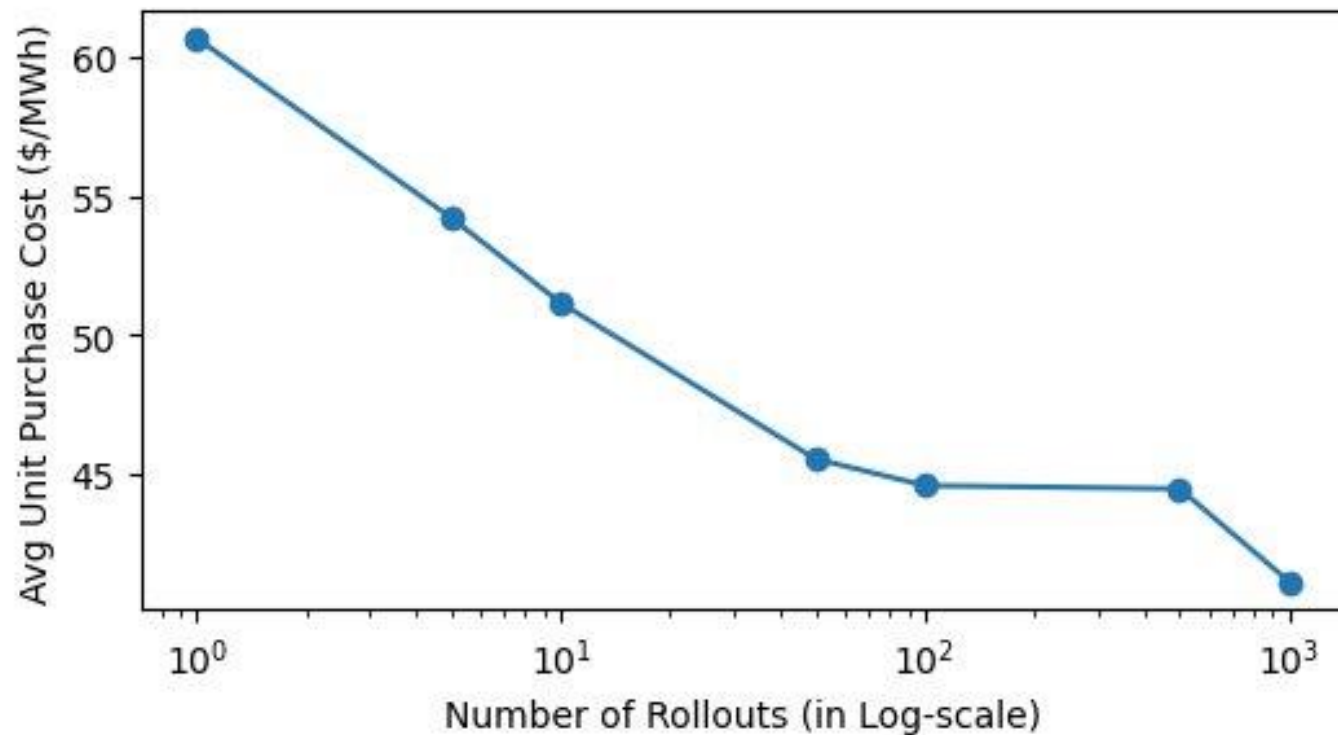
PDA Bidding Strategies:

ZI: Strategy to place random bids within upper and lower bounds

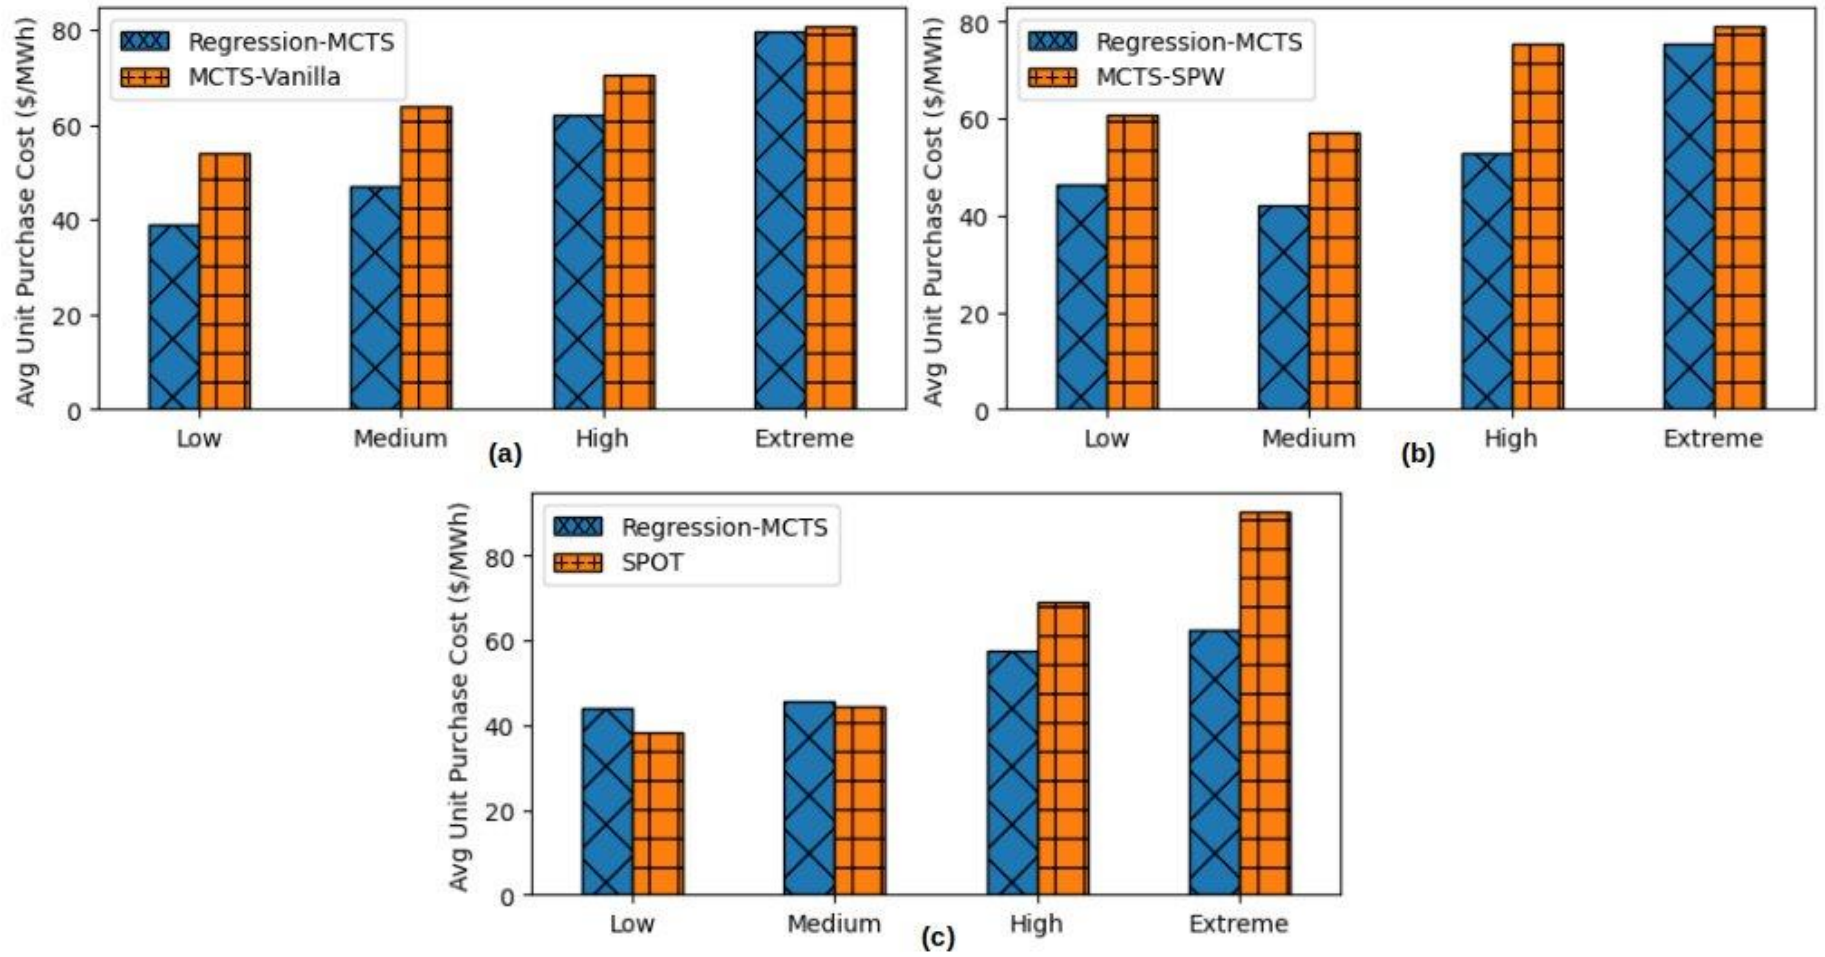
ZIP: Heuristic strategy to adjust the profit margins while placing bids

VV21: Heuristic strategy that models the cost curve of the generating company derived from uncleared ask information available in PDAs. Then, uses the cost curve information to place bids.

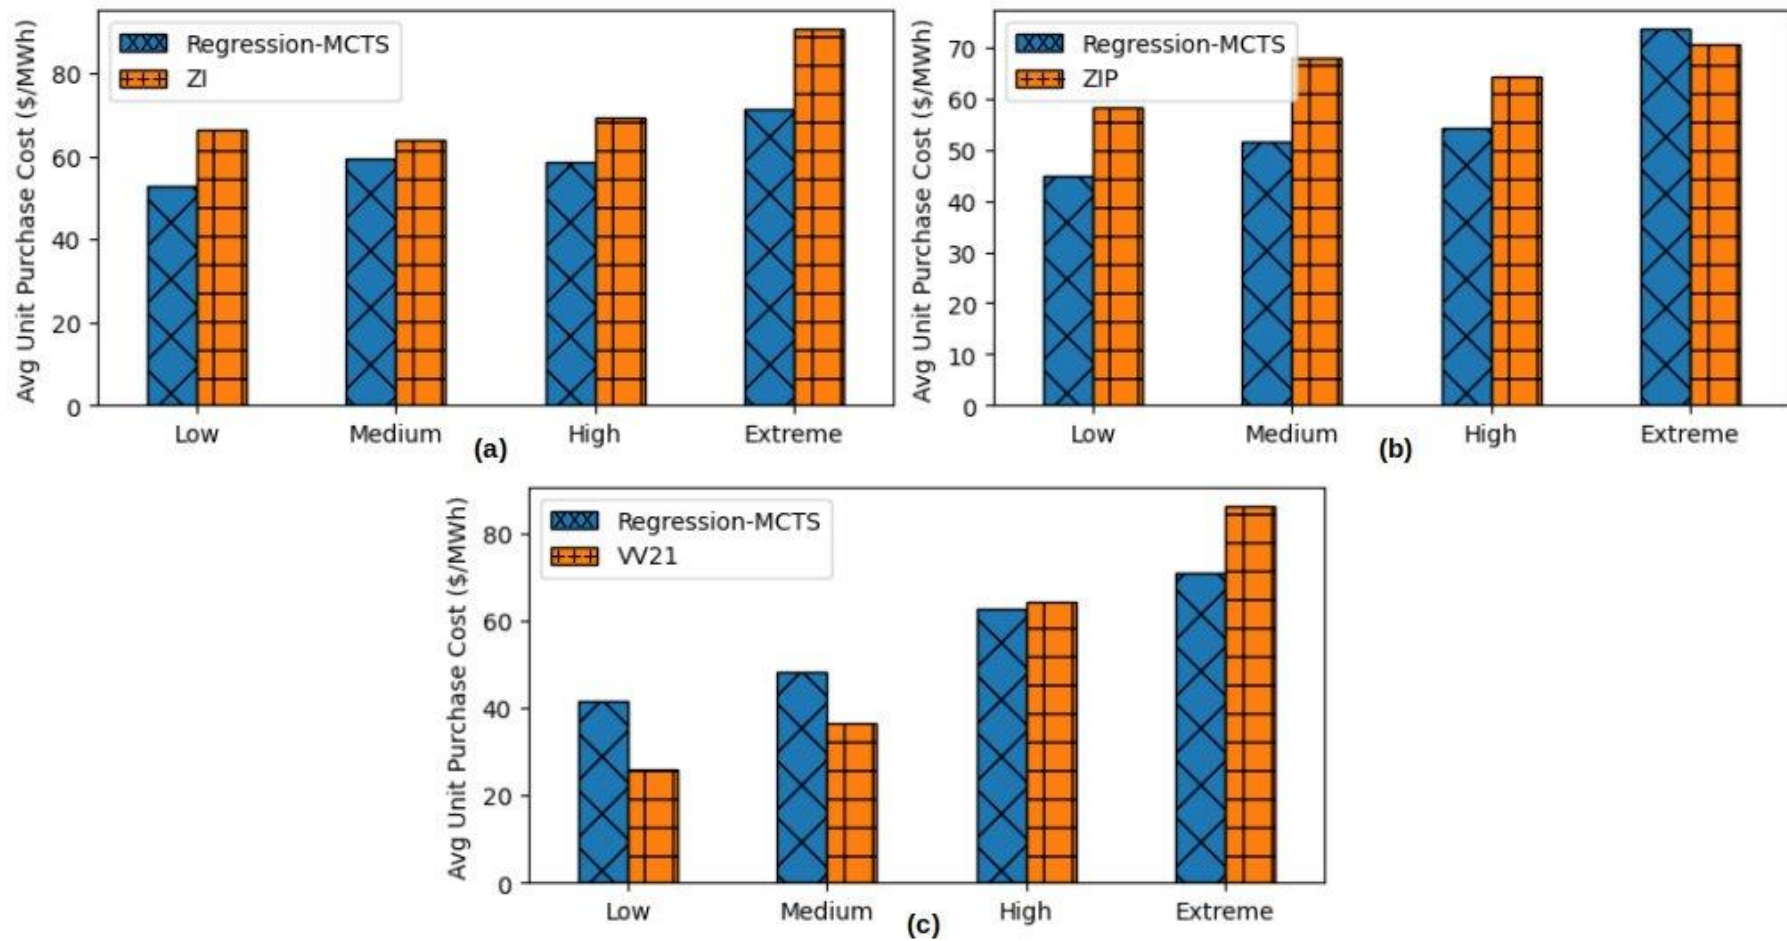
Set-1: Performance R-MCTS with #rollouts



Set-2: Performance R-MCTS against state-of-the-art MCTS methods



Set-3: Performance R-MCTS against PDA bidding strategies



Thank you !!!

Questions ... ?